The background of the slide features a photograph of a dense orchard. The trees are bare or have sparse, light-colored blossoms. In the foreground, there is a grassy area with small yellow flowers, possibly dandelions. The overall tone is soft and slightly overexposed.

S&DS 265 / 565
Introductory Machine Learning

Trees and Random Forests

October 6

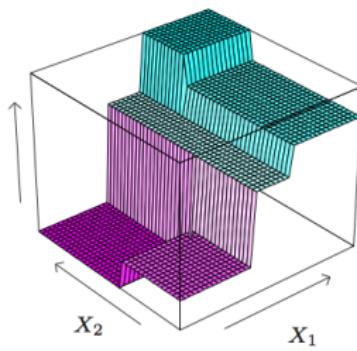
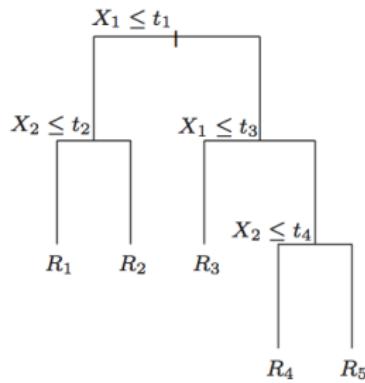
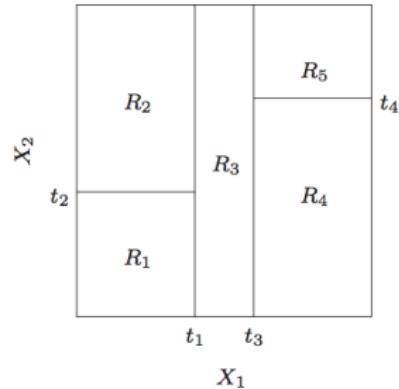
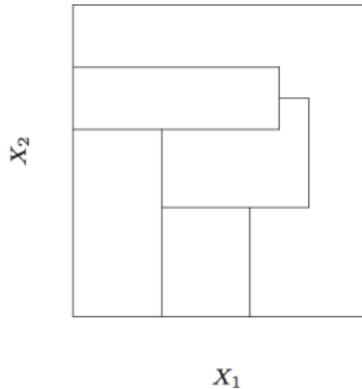
Yale

Reminders

- Assn 2 out; due week from today
- Quiz 3 available 10:30am today (48 hours, 20 minutes)
 - ▶ Classification concepts, bias-variance tradeoff, SGD, trees
- Midterm in class on October 18
- Practice midterms have been posted
- Review sessions TBA
- Questions?

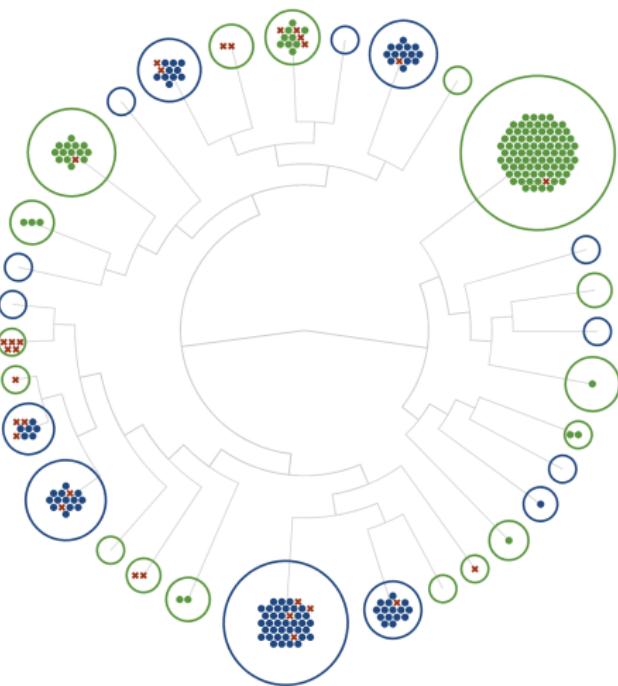
Last time: trees

- Give interpretable decision rules
- Deep trees have low bias, high variance
- Trees can be grown greedily to be full, then pruned back
- Predictive power is ... meh



Beautiful demo

<http://www.r2d3.us/visual-intro-to-machine-learning-part-2>



Classification Trees

- \hat{y}_i for all $i \in R_j$ is *most commonly occurring class* of training observations in R_j .

$$\hat{y}_{R_j} = \arg \max_k \hat{p}_{jk},$$

where \hat{p}_{jk} is the proportion of training observations with label k in rectangle (leaf) R_j .

- No longer want to minimize RSS, but instead minimize...
 - ▶ classification error rate

$$E = \sum_{j=1}^J |R_j| (1 - \max_k (\hat{p}_{jk}))$$

Classification Trees

- \hat{y}_i for all $i \in R_j$ is *most commonly occurring class* of training observations in R_j .

$$\hat{y}_{R_j} = \arg \max_k \hat{p}_{jk},$$

where \hat{p}_{jk} is the proportion of training observations with label k in rectangle (leaf) R_j .

- No longer want to minimize RSS, but instead minimize...
 - ▶ classification error rate

$$E = \sum_{j=1}^J |R_j| (1 - \max_k (\hat{p}_{jk}))$$

- ▶ Gini index

$$G = \sum_{j=1}^J |R_j| \sum_{k=1}^K \hat{p}_{jk} (1 - \hat{p}_{jk})$$

Impurity measures

Define node proportion of class k

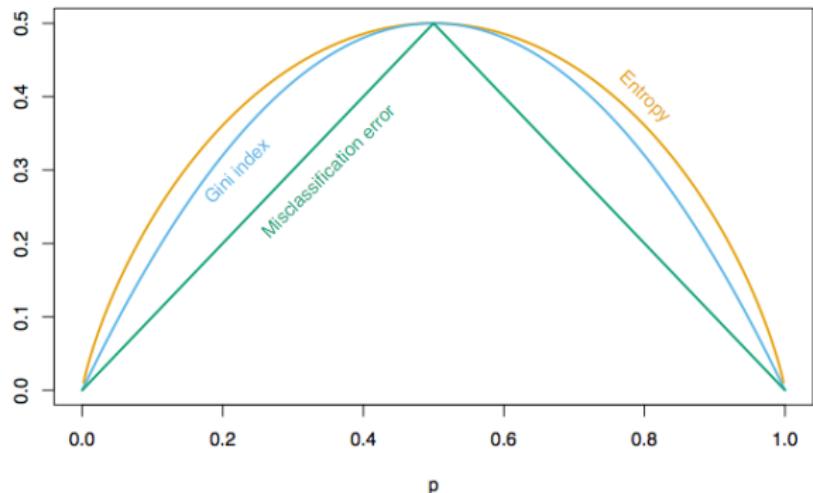
$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \in R_m} I(y_i = k)$$

$$k(m) = \arg \max_k \hat{p}_{mk}$$

- Misclassification error: $1 - \hat{p}_{mk(m)}$
- Gini index: $\sum_{k=1}^K \hat{p}_{mk} (1 - \hat{p}_{mk})$
- Entropy: $-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$

*All of these encourage higher **node purity***

Impurity measures



Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

- Both produce simple predictions (averages/maximally occurring) based on “neighborhoods” in the predictor space.

Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

- Both produce simple predictions (averages/maximally occurring) based on “neighborhoods” in the predictor space.
- Neighborhoods chosen very differently

Trees vs. other methods

Recall that linear regression fits models of the form

$$f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$$

Trees vs. other methods

Recall that linear regression fits models of the form

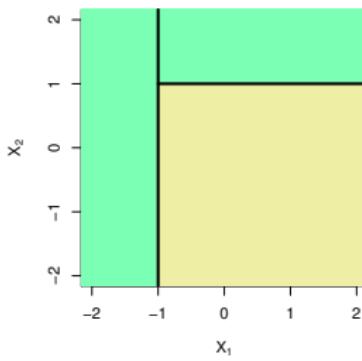
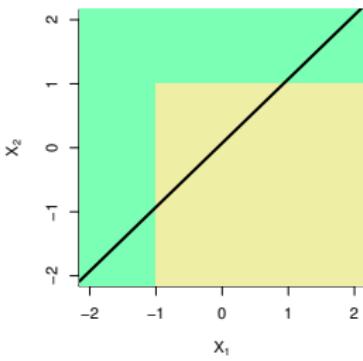
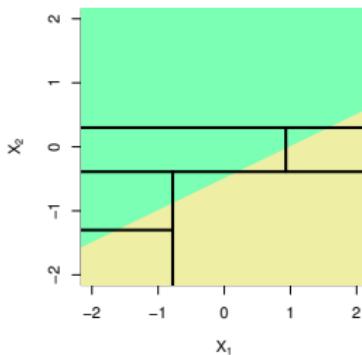
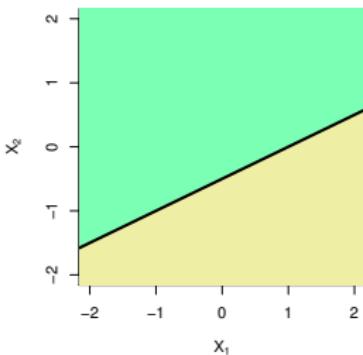
$$f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$$

Regression trees are like fitting linear regression models with a bunch of indicators

$$f(X) = \sum_{j=1}^J \beta_j \mathbb{1} \{X \in R_j\}$$

Trees vs. other methods

Are trees always better than linear methods?



Summary

- trees are intuitive
- prediction rules easy to explain, interpret
- trees are sensitive to underlying data
- trees produce non-smooth prediction surfaces
- prediction accuracy can be so-so

Ensemble methods

Ensemble methods pool together multiple models to arrive at more reliable predictions.

Ensemble methods

Ensemble methods pool together multiple models to arrive at more reliable predictions.

- bagging
- random forests
- boosting

Bootstrap samples

- We have n data points
- Randomly sample exactly n of them *with replacement*
- Contains about $\frac{2}{3}$ of the data, with duplicates
- This is a “bag” of n data points

Chance a particular example does not appear in sample:

$$\left(1 - \frac{1}{n}\right)^n \longrightarrow \frac{1}{e} \approx 0.37$$

Bagging

Regression trees: If we had multiple training sets, could grow multiple trees, then take an average.

Bagging

Regression trees: If we had multiple training sets, could grow multiple trees, then take an average.

Create B bootstrap samples, grow tree (without pruning) using each $\hat{f}^{*1}, \hat{f}^{*2}, \dots, \hat{f}^{*B}$. For prediction at x , we take an average:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

Bagging

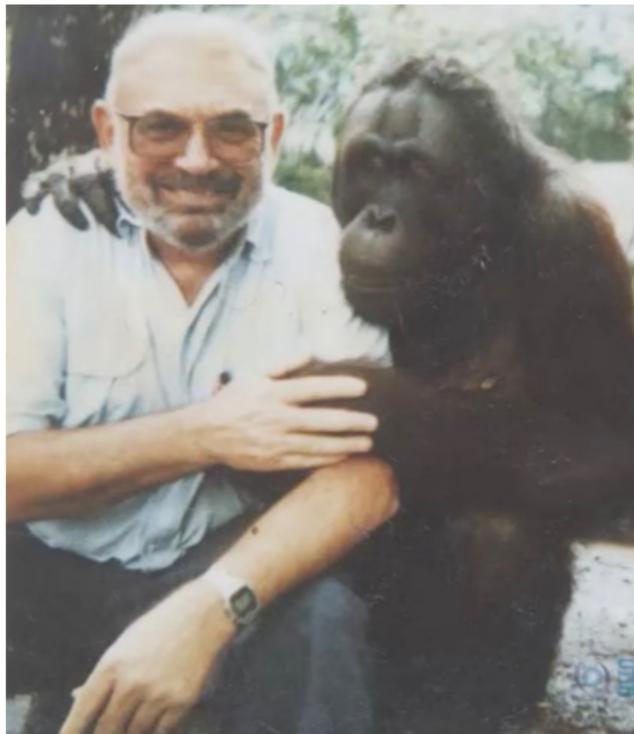
Regression trees: If we had multiple training sets, could grow multiple trees, then take an average.

Create B bootstrap samples, grow tree (without pruning) using each $\hat{f}^{*1}, \hat{f}^{*2}, \dots, \hat{f}^{*B}$. For prediction at x , we take an average:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

Classification trees: For prediction at x , $\hat{f}_{bag}(x)$ is decided by *majority vote*.

Leo Brieman – “Keep it Simple”



Out-of-Bag error estimation

Each bagged tree uses about 2/3 of all observations (with repeats).

The remaining data—*out-of-bag* (OOB) observations—can be put to good use.

Out-of-Bag error estimation

Obs	Bagging iteration					OOB Est.
	1	2	3	...	B	
1	OOB	train	train	...	train	\hat{y}_1
2	train	OOB	train	...	train	\hat{y}_2
3	train	train	OOB	...	train	\hat{y}_3
4	OOB	train	train	...	OOB	\hat{y}_4
...
n	train	train	OOB	...	train	\hat{y}_n

Out-of-Bag error estimation

Obs	Bagging iteration					OOB Est.
	1	2	3	...	B	
1	OOB	train	train	...	train	\hat{y}_1
2	train	OOB	train	...	train	\hat{y}_2
3	train	train	OOB	...	train	\hat{y}_3
4	OOB	train	train	...	OOB	\hat{y}_4
...
n	train	train	OOB	...	train	\hat{y}_n

- For each bagged tree, we can make predictions for the OOB observations.
- At the end, we can aggregate over all predictions for the i -th observation to arrive at a OOB prediction \hat{y}_i .
- We can compute prediction error based on these OOB predictions $\hat{y}_1, \dots, \hat{y}_n$.

Out-of-Bag error estimation

Obs	Bagging iteration					OOB Est.
	1	2	3	...	B	
1	OOB	train	train	...	train	\hat{y}_1
2	train	OOB	train	...	train	\hat{y}_2
3	train	train	OOB	...	train	\hat{y}_3
4	OOB	train	train	...	OOB	\hat{y}_4
...
n	train	train	OOB	...	train	\hat{y}_n

- For each bagged tree, we can make predictions for the OOB observations.
- At the end, we can aggregate over all predictions for the i -th observation to arrive at a OOB prediction \hat{y}_i .
- We can compute prediction error based on these OOB predictions $\hat{y}_1, \dots, \hat{y}_n$.

Variable importance

While bagging improves upon the predictive ability of trees, it kills off the interpretability of the model.

Variable importance

While bagging improves upon the predictive ability of trees, it kills off the interpretability of the model.

A good tool for interpreting a bagged tree model (and other tree-based ensemble methods like forests) is the **variable importance measure**.

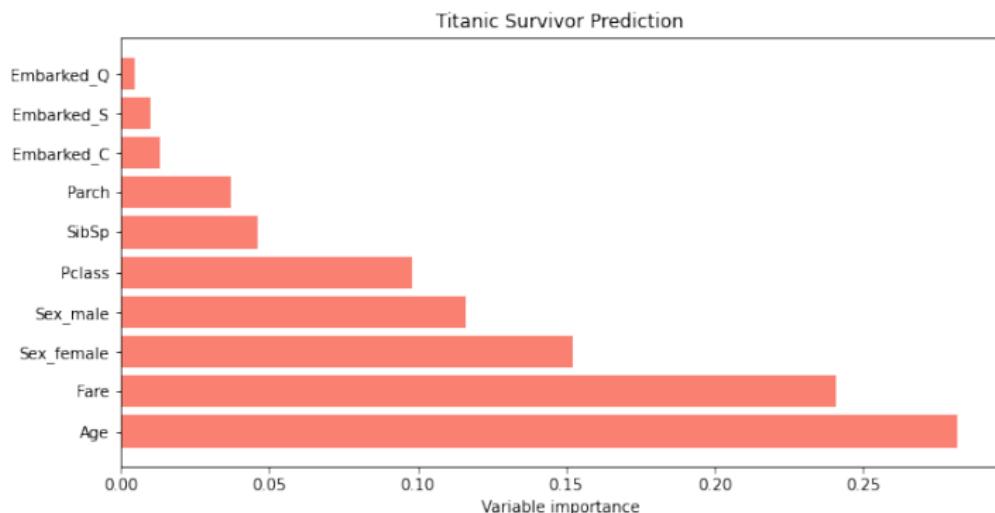
Variable importance

While bagging improves upon the predictive ability of trees, it kills off the interpretability of the model.

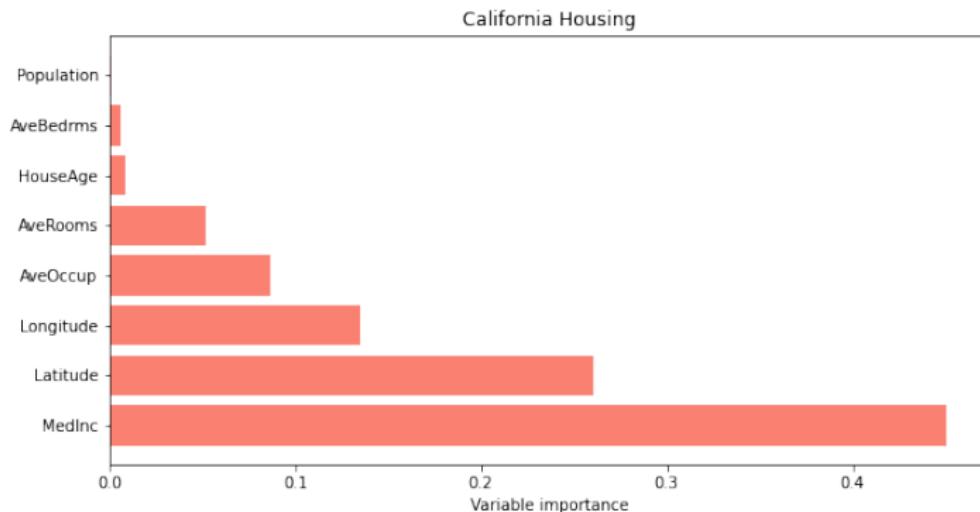
A good tool for interpreting a bagged tree model (and other tree-based ensemble methods like forests) is the **variable importance measure**.

Variable importance can be measured by the amount that the RSS (or Gini index) is reduced due to splits over a given predictor, averaged over all B trees.

Variable importance



Variable importance



Random Forests

Similar to bagging, but takes the averaging idea even further –
averaging uncorrelated things decreases the error!

Random Forests

Similar to bagging, but takes the averaging idea even further –
averaging uncorrelated things decreases the error!

Still use bootstrap samples, but only use m out of p predictors at
each split, *chosen randomly*

Bagging and Random Forests

- Grow many trees and average their predictions
- Trees are grown deep, to have low bias, but high variance
- To “decorrelate” the predictors, each tree is
 - ▶ grown on a bootstrap sample of the data
 - ▶ grown with random subsets of the predictors at each split
- Tree growing can be done in parallel

Guess who?



Random Forests Algorithm

- ① For $b = 1$ to B :
 - (a) Draw a bootstrap sample Z^* of size n from the training data
 - (b) Grow a random-forest tree T_b to the bootstrapped data, recursively repeating following steps, until minimum node size reached:
 - i. Select m variables at random from the p variables
 - ii. Pick the best variable/split-point among the m
 - iii. Split the node into two children nodes
- ② Output the ensemble of trees $\{T_b\}_{b=1}^B$.

To make a prediction at a new point x :

Regression: Average $\widehat{f}_{rf}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$

Classification: Majority vote of the individual trees

Random forests–recommended parameters

- For classification, default value of m is $\lfloor \sqrt{p} \rfloor$ and the minimum node size is one.
- For regression, the default values of m is $\lfloor p/3 \rfloor$ and the minimum node size is five.

(you don't need to remember this)

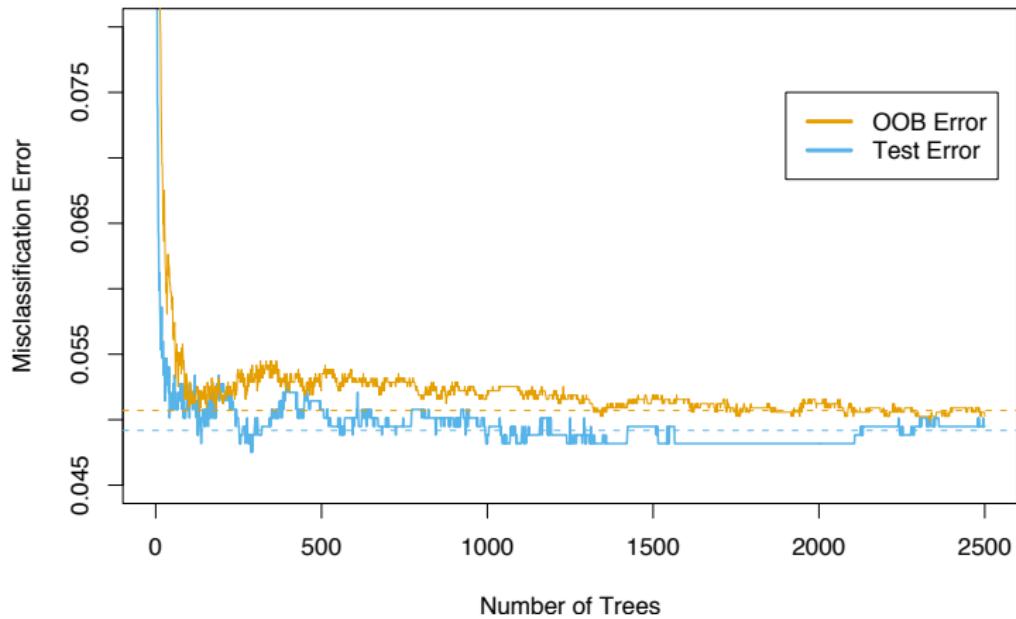
Out of bag (OOB) prediction

As before, can use out-of-bag (OOB) samples:

- For each observation $z_i = (x_i, y_i)$, construct its random forest predictor by averaging only those trees corresponding to bootstrap samples in which z_i did not appear
- Thus, cross-validation can be performed “along the way”

Chance a sample x_i does not appear in a bootstrap sample is $\left(1 - \frac{1}{n}\right)^n \longrightarrow \frac{1}{e} \approx 0.37$

Out of bag (OOB) prediction



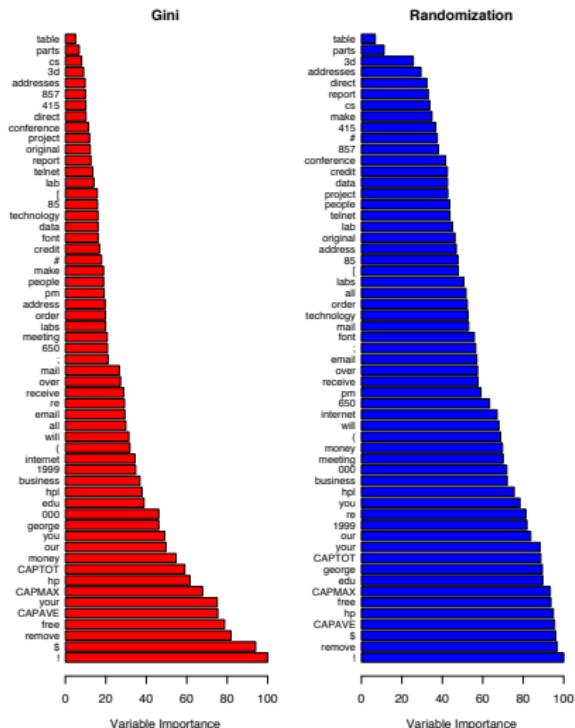
Performance on email spam task

Single tree: 8.7%

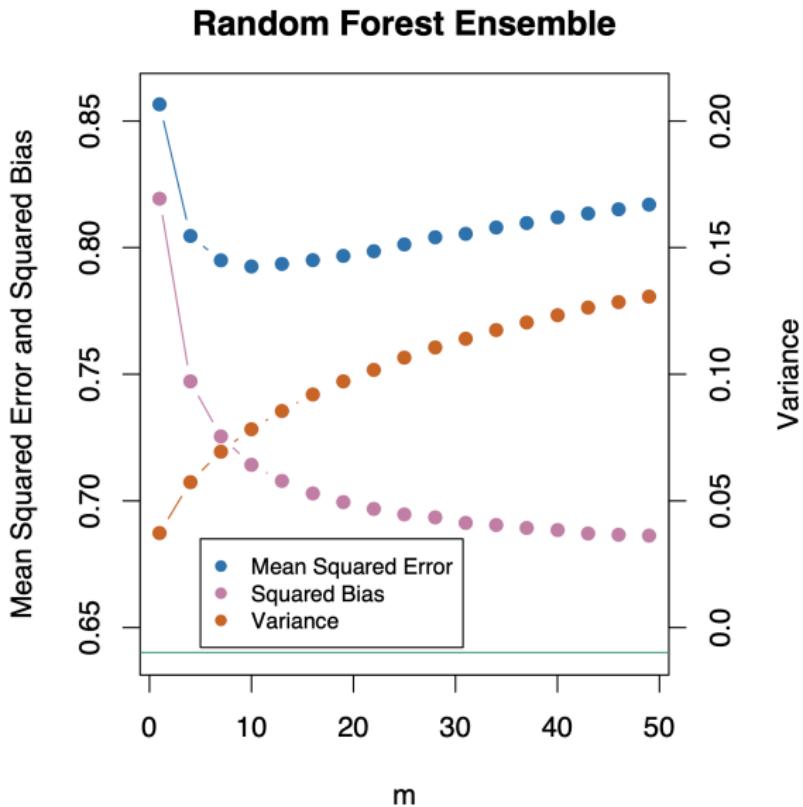
Random forests: 5.1%

(standard error of the estimates is $\approx 0.6\%$)

Important features: Spam



Random forest MSE



Let's go to the notebook

```
In [1]: import numpy as np  
import matplotlib.pyplot as plt  
import pandas as pd
```

```
In [2]: titanic_train = pd.read_csv('https://raw.githubusercontent.com/minsuk-heo/kaggle-titanic/master/input/train.csv')  
titanic_test = pd.read_csv('https://raw.githubusercontent.com/minsuk-heo/kaggle-titanic/master/input/test.csv')  
titanic_train
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1		female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W.C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q

891 rows x 12 columns

What did we learn today?

- Bagging tries to reduce variance by averaging many trees
- Random forests decorrelate by random sampling of predictors
- Try to manage bias-variance tradeoff with randomization